



Week 3 Virtualization and Containerization

Dr Zhi Zhang

Overview

- Operating System Concepts
- Virtualization
- Containerization

Operating System Concepts

- Process scheduling
- Virtual memory

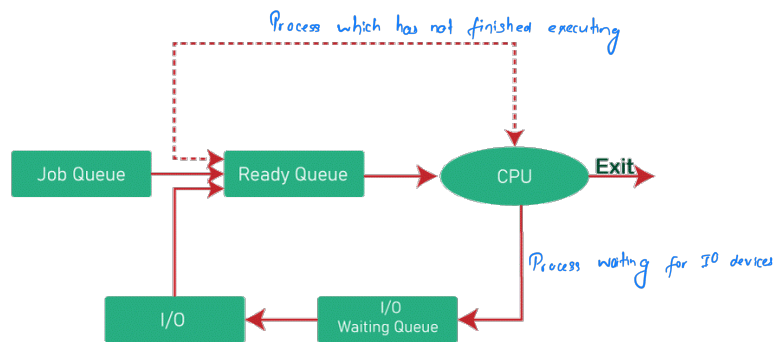
What is a process?

- A process is a program that is executed
- It is a basic unit of execution in an OS.
- At a minimum, process execution requires **CPU and memory without I/O devices.**

Process scheduling (time-sharing)


- Each process is scheduled and executed within their time share.
- The number of processes can be much larger than that of the CPUs.

Process scheduling



<https://www.tutorialandexample.com/process-scheduling-in-operating-system>


Virtual memory



https://www.cs.uic.edu/~jbell/CourseNotes/OperatingSystems/9_VirtualMemory.html

Virtual memory

- **Page table:**
 - In a modern OS (e.g., Linux, Windows or MacOS), the OS uses a set of page tables to map virtual memory within a process to their corresponding physical memory in main memory.

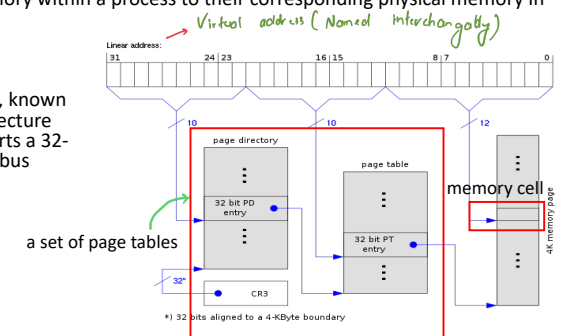


- VA refers to virtual address, pointing to a memory cell in a **virtual page**.
- PA refers to physical address, pointing to a memory cell in a **physical page**.

https://www.alibabacloud.com/blog/a-tribute-to-hackers-the-way-to-explore-memory-virtualization_599058

Virtual memory

- **Page table:**
 - In a modern OS (e.g., Linux, Windows or OS X), the OS uses a set of page tables to map virtual memory within a process to their corresponding physical memory in main memory.



<https://www.cs.virginia.edu/~bjc8c/class/cs6456-s20/slides/01-intro.pptx>

What is virtualization?

- Virtualization is the ability to run multiple operating systems on a single physical system and share the underlying hardware resources [1]
- Allows one computer software (called **Virtual Machine Monitors** or **Hypervisor**) to provide the appearance of many computers (called virtual machines).
- Goals:
 - Provide flexibility for users
 - Amortize hardware costs
 - Isolate completely separate users

[1] VMWare white paper, Virtualization Overview

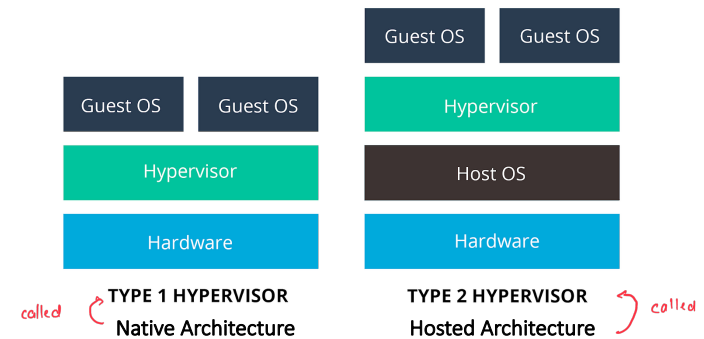
Characteristics for Virtualizable Computer Architectures

- First, the VMM provides an environment for VM. The environment is essentially identical with the original machine.
- Second, VM running in this environment shows minor decreases in speed.
- Last, the VMM is in complete control of all system resources, such as CPUs, memory, disk, etc.

VMM Types

- **Hosted Architecture**
 - Install as an application on existing x86 “host” OS, e.g. Windows, Linux, OS X
 - Leverage host I/O stack and resource management
 - Examples: Virtualbox, VMware
- **Bare-Metal/Native Architecture**
 - VMM or Hypervisor is installed directly on hardware
 - Acknowledged as preferred architecture for mainstream public clouds
 - Examples: KVM, Xen, Hyper-V

VMM Types



Virtualization: history

- 1960's: first track of virtualization
 - Time and resource sharing on expensive computers
 - IBM VM/370
- Late 1970s and early 1980s: became unpopular
 - Cheap hardware and multiprocessing OS
- **Late 1990s:** became popular again
 - Wide variety of OS and hardware configurations (compatibility issue)
 - VMWare
- **Since 2000:** hot and important
 - Cloud computing
 - Containers

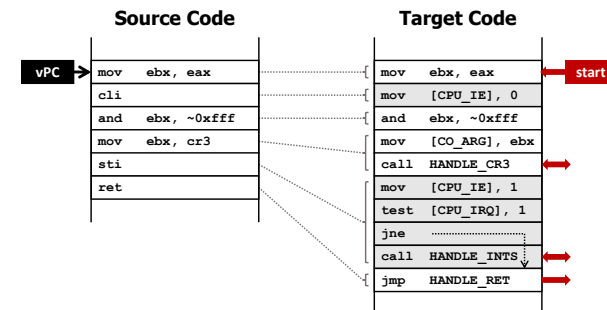
Virtualization on x86 architecture

- Solutions: → *Generations of virtualization solutions*
 - Full virtualization implemented by Dynamic binary translation & Two-stage translation
 - Para-virtualization (Xen)
 - Full virtualization implemented by Hardware extension

Dynamic Binary Translation

- The entire program (e.g., ARM format) is translated into a binary of another architecture (e.g., x86 format).
- The translation occurs on-the-fly during program execution.

Dynamic Binary Translation




Courtesy Scott DeWine VMware Inc

Dynamic Binary Translation

- For sensitive/privileged instructions, trap-and-emulate is needed.
 - e.g., IBM VM/370: the first virtualized system


IBM VM/370: the first virtualized system



<https://www.cs.virginia.edu/~bjc8c/class/cs6456-s20/slides/01-intro.pptx>

IBM VM/370

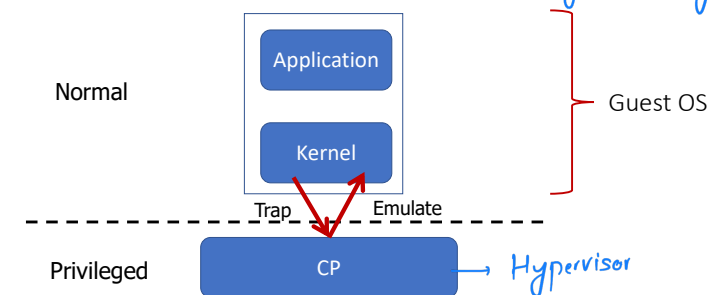
- Virtualization technology: **trap-and-emulate**



This is how normal OS would look. Forget about Guest OS here.

<https://www.cs.virginia.edu/~bjc8c/class/cs6456-s20/slides/01-intro.pptx>

- When guest OS wants to run some privileged instructions, those instructions will create a trap that goes into hypervisor and hypervisor will emulate the intended functionality of the guest OS*
- Virtualization technology: **trap-and-emulate**




<https://www.cs.virginia.edu/~bjc8c/class/cs6456-s20/slides/01-intro.pptx>

Two-stage translation

- **Guest page table:**
 - A virtual machine has its own guest page table (GPT), managed by the guest OS. The GPT translates guest virtual addresses to guest physical addresses.
- **Shadow page table:**
 - The hypervisor creates a shadow page table for each guest page table. The shadow page table translates guest virtual addresses to machine addresses.
- **Two-stage translation:**
 - When a virtual machine running on the hypervisor performs a memory access, the virtual address goes through two stages of translation.

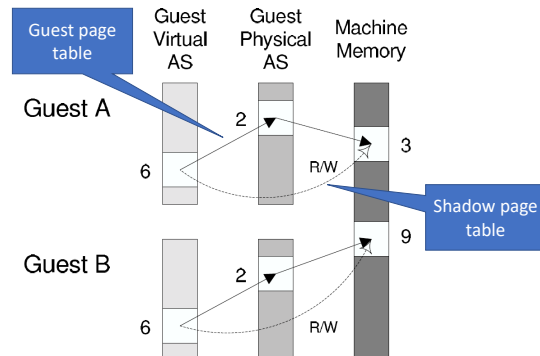


Two-stage translation



<https://www.cs.virginia.edu/~bjc8c/class/cs6456-s20/slides/01-intro.pptx>

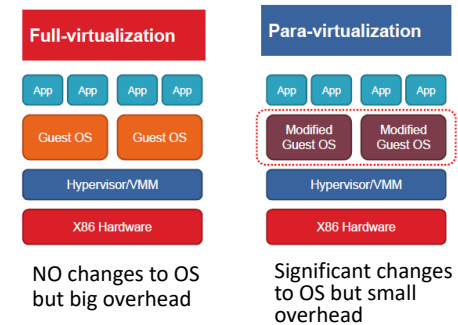
Two-stage translation



<https://www.cs.virginia.edu/~bjc8c/class/cs6456-s20/slides/01-intro.pptx>


Para-virtualization

- Full vs. Para virtualization



A typical example: Xen


- SOSP (ACM Symposium on Operating Systems Principles) in 2003.
- Very high impact (data collected in 2013 and 2023)



Overview of Xen

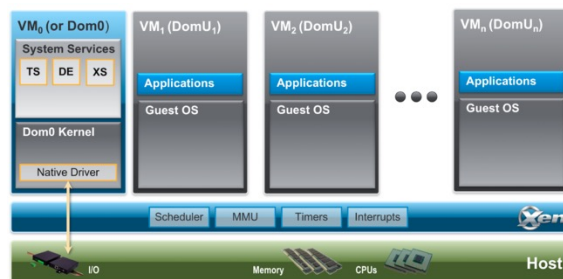
- Modify guest OS to be aware of virtualization
 - No emulation of privileged operations from guest OS
 - Better performance

(a) MMU Para-virtualization



Single Translation stage but compatibility issue

Xen architecture



- TS is Toolstack, providing a user interface to the Xen hypervisor administrator.
- DE is Device Emulation (DE), emulating devices when the devices are not available.
- XS is XenStore/XenBus, managing information that are shared between some domains.

https://wiki.xenproject.org/wiki/Xen_Project_Software_Overview

A short conclusion


- Full virtualization
 - Unmodified guest OS
 - Performance issue
- Para virtualization:
 - Better performance
 - Modified guest OS

Support from hardware extensions

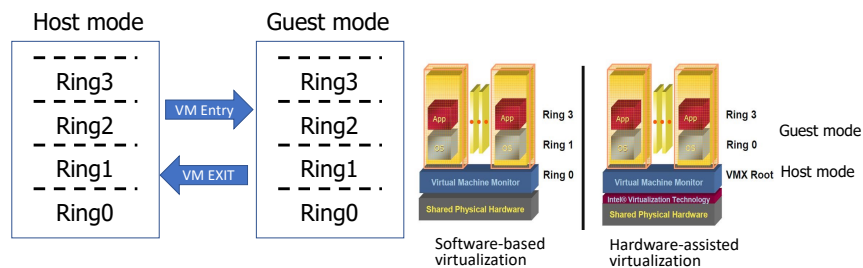
- Full/Para virtualization: **software-based** virtualization
- **Hardware support**: Intel and AMD assist virtualization

x86 Protection Rings (CPL)

- x86 has four protection levels (rings) /CPU privilege levels (CPLs).
 - ring 0 – “Kernel mode” (most privileged)
 - ring 3 – “User mode”
 - ring 1 & 2 – Other
- Linux only uses ring 0 and ring 3.
 - “Kernel vs. user mode”
- For previous full virtualization, guest OS applications run in ring 3, guest OS kernel runs in ring 1, and the hypervisor runs in ring 0.




CPU virtualization: Intel VT-x & AMD SVM



An example: Intel VT-x

- Two VT-x operating modes
 - Less-privileged mode (guest or VMX non-root) for guest OSes
 - More-privileged mode (host or VMX root) for VMM
- Two transitions
 - VM entry to non-root mode
 - VM exit to root mode



CPU virtualization

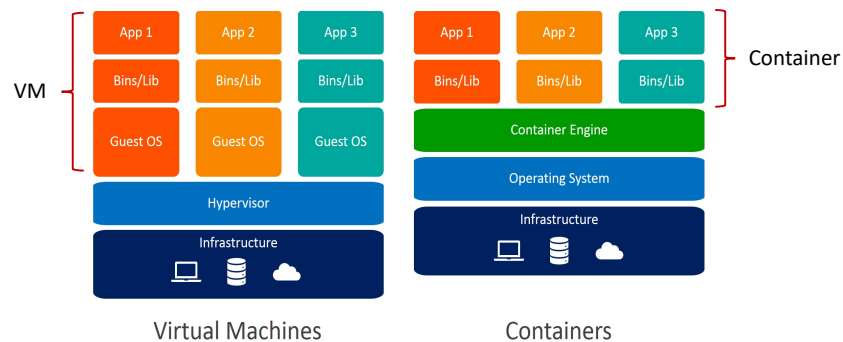
- New VM mode bit
 - Create two orthogonal modes
- If VM mode bit is **clear** → host mode
- If VM mode bit is **set** → guest mode
- Benefits:
 - Eliminating the need of binary translation: significant performance improvement.
 - Requiring no changes to privileged operations in guest OSes: compatible with existing modern OSes.

Containers: concepts

- A container is a sandboxed process running on a host OS that is isolated from all other processes running on that host OS [2]
- A container image is a stand-alone and executable software package (e.g., dependencies, binaries, etc) that contains everything needed to run an application.
- A container is a runnable instance of the image.

[2] <https://docs.docker.com/get-started/>

Containers



<https://www.weave.works/blog/a-practical-guide-to-choosing-between-docker-containers-and-vm>

Containerization vs. Virtualization


- Granularity
 - Containers are an abstraction of the process layer and VMs are a simulation of the hardware layer.
- Overhead
 - Required resources: containers are created to run one application and VMs support a whole OS.
 - Efficiency: containers are launched to run an application. VMs need to boot up an entire OS.
- Security/Isolation
 - Containers are isolated from each other at the process level. VMs are isolated at the OS level.

Backbones of containers

- namespaces
- cgroups

namespaces: limit what a container can see


- Provide containers with their own view of system resources (e.g., network)



<https://subscription.packtpub.com/book/cloud-and-networking/9781838827472/13/ch13v1sec97/running-in-an-existing-network-namespace>


cgroups: limits how much a container can use

- Control Groups: manage resource usage and limits for processes within a container.
 - Example: a container is limited to use only 1 Mbps of network bandwidth



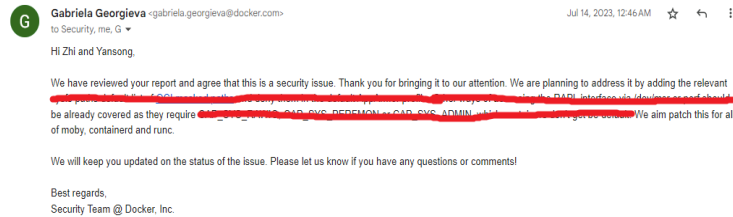
Docker container

- Docker provides an interface on top of the techniques
 - Popularized containers (a standard unit of software)



<https://medium.com/@saschagrunert/demystifying-containers-part-i-kernel-space-2c53d6979504>

Docker container



Docker cheatsheet

- `docker run` – Runs a command in a new container.
- `docker start` – Starts one or more stopped containers
- `docker stop` – Stops one or more running containers
- `docker build` – Builds an image from a Docker file
- `docker pull` – Pulls an image or a repository from a registry
- `docker push` – Pushes an image or a repository to a registry
- `docker export` – Exports a container's filesystem as a tar archive
- `docker exec` – Runs a command in a run-time container
- `docker search` – Searches the Docker Hub for images
- `docker attach` – Attaches to a running container
- `docker commit` – Creates a new image from a container's changes

Practice Questions

- [7 marks] Q1: Describe what virtualization is and describe the characteristic attributes of the different types of virtualization (Language, Operating System and Hardware).
- [1 mark] Virtualization is a technology that allows multiple virtual instances of resources (e.g., operating systems) to run on a single physical machine.
- [2 marks] **Language virtualization:** provides a virtual runtime environment (e.g., JVM) that allows code written in a specific programming language (e.g., Java) to execute.
 - Cross-OS: code can be executed in a runtime environment across OSes without modification.
- [2 marks] **Operating system virtualization:** allows multiple user-space instances (e.g., docker containers) within a single operating-system kernel (e.g., Linux kernel).
 - Limited Isolation: instances share the host OS kernel, which means a vulnerability in the kernel could affect all instances.
- [2 marks] **Hardware virtualization:** creates multiple virtual machines sharing a single physical server.
 - VM Snapshots: allows saving the current state of a VM
 - Live VM Migration: enables moving runtime VMs between physical hosts seamlessly.