

HELP US ALL STAY HEALTHY

SIX SIMPLE TIPS



Maintain 1.5 metres distance
between yourself and others
where possible



Cough and sneeze into
your elbow or a tissue
(not your hands)



Avoid shaking hands



Put used tissues
in the bin



Wash hands with soap and
warm water or use an alcohol-
based hand sanitiser after you
cough or sneeze



Do not touch
your face

IF YOU ARE UNWELL AND WORRIED ABOUT COVID-19:

- Call the National
Coronavirus Helpline:
1800 020 080
- Call your usual GP for advice
- Call the UWA Medical Centre
for advice: 6488 2118

UWA FAQs:
uwa.edu.au/coronavirus

Report COVID-19 hazards
and suspected/confirmed
cases via RiskWare:
uwa.edu.au/riskware



THE UNIVERSITY OF
**WESTERN
AUSTRALIA**

High-Performance Computing

Lecture 4 OpenMP Tasking and Synchronisation

CITS5507

Zeyi Wen

Computer Science and
Software Engineering

School of Maths, Physics
and Computing

Acknowledgement: The lecture slides are adapted from many online sources.

- **OpenMP Task**
 - ✓ What is task
 - ✓ Task generation and execution
 - ✓ Data environment of tasks
 - ✓ Task Switching
- Synchronisation
 - ✓ Why synchronisation
 - ✓ Synchronisation constructs
 - ✓ Critical and atomic
 - ✓ Flush and lock

- Key concept: OpenMP always had tasks, the parallel computing community just never called them that.
 - Thread encountering *parallel construct* packages up a set of implicit tasks, one per thread.
 - Team of threads is created. Each thread in team is assigned to one of the tasks (and tied to it).
 - Barrier holds original master thread until all implicit tasks are finished.

```
#pragma omp parallel
{
    #pragma omp for
    for(i=0; i<N; i++){
        c[i] = a[i]+b[i];
    }
}
```

Sequential code

```
for(i=0; i<N; i++){
    c[i] = a[i]+b[i];
}
```

Assume used 4 threads

(Lecture 3) Loop Worksharing Construct

- The loop worksharing construct splits up loop iterations among the threads in a team.

```
#pragma omp parallel
{
    #pragma omp for
    for(i=0; i<N; i++){
        do_something(i);
    }
}
```

```
#pragma omp parallel
{
    #pragma omp for
    for(i=0; i<N; i++){
        c[i] = a[i]+b[i];
    }
}
```


Sequential code

```
for(i=0; i<N; i++){
    c[i] = a[i]+b[i];
}
```

Assume used 4 threads

(Lecture 3) Loop Worksharing Example

Vector Addition: $\mathbf{c} = \mathbf{a} + \mathbf{b}$



- Key concept: OpenMP always had tasks, the parallel computing community just never called them that.
 - Thread encountering *parallel construct* packages up a set of implicit tasks, one per thread.
 - Team of threads is created. Each thread in team is assigned to one of the tasks (and tied to it).
 - Barrier holds original master thread until all implicit tasks are finished.
- OpenMP 3.0 has simply added a way to create a task explicitly for the team to execute.
- Every part of an OpenMP program is part of one task or another!

The official introduction of the **task** construct:

- The **task** construct defines an **explicit task**, which may be executed by the encountering thread or **deferred** for execution by any other thread in the team.
- The **data environment** of the task is determined by the **data sharing attribute clauses**.
- Task execution is subject to task scheduling.
 - ✓ Also see the OpenMP 3.1 documentation for the associated **taskyield** and **taskwait** directives.

The syntax of task construct

```
#pragma omp task [clause [ [,] clause] ...]  
    //structured code block
```

where **clause** can be one of:

- if (expression)
- untied
- shared(list)
- private(list)
- firstprivate(list)
- default(shared | none)

We will
talk later...

- **OpenMP Task**
 - ✓ What is task
 - ✓ Task generation and execution
 - ✓ Data environment of tasks
 - ✓ Task Switching
- Synchronisation
 - ✓ Why synchronisation
 - ✓ Synchronisation constructs
 - ✓ Critical and atomic
 - ✓ Flush and lock

Example: how many tasks generated, and why?

```
/* Create threads */  
#pragma omp parallel num_threads(2)  
{  
    #pragma omp task  
        t0();  
    #pragma omp task  
        t1();  
}  
// Implicit barrier
```

4, as two threads
generate two tasks
respectively.

(Lecture 3) Single Worksharing Construct

- The **single** construct denotes a block of code that is executed by only one thread (not necessarily the master thread).
- A barrier is implied at the end of the single block (can remove the barrier with a **nowait** clause).

```
#pragma omp parallel
{
    do_many_things();

#pragma omp single
{
    exchange_boundaries();
}

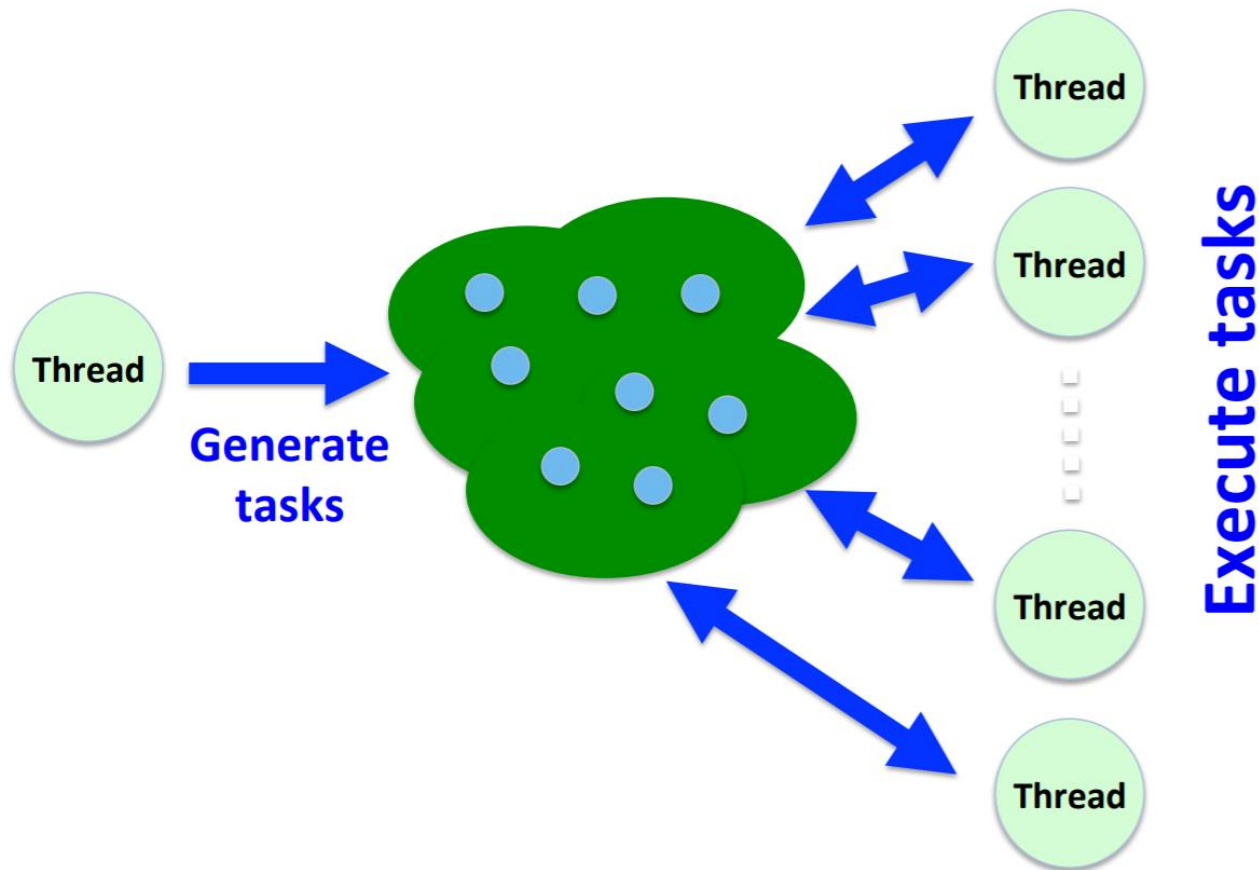
    do_many_other_things();
}
```

The typical skeleton of a **task** construct

```
/* Create threads */
#pragma omp parallel num_threads(2)
{
    #pragma omp single
    {
        #pragma omp task
        t0();
        #pragma omp task
        t1();
    }
}
// Implicit barrier
```

One generator,
multiple tasks.

Task Generation and Execution



“For” and “sections” directive defects :

- Traditional implicit tasks within “for” and “section” constructs **cannot** be **dynamically** partitioned based on the **runtime environment**;
- The schedule of tasks in “for” and “sections” **must be known in advance** and the “**task**” directive are used to solve this problem.

- When a thread encounters a **task** construct, it may choose to execute the task **immediately** or **defer** execution until some later time.
- If a task is delayed, it is placed in the **conceptual task pool** associated with the current parallel region. Threads in the current group **pull tasks** out of the pool and execute them until the pool is empty.

(The thread executing the task may be different from the thread that initially encountered the task.)

Example: the defects of “for”

```
int main()
{
    int a[N];
    init(a, N);
    #pragma omp parallel num_threads(2)
    {
        #pragma omp for
        for(int i = 0; i < N; i++)
        {
            process(a[i]);
        }
    }
    return 0;
}
```

The “for” directive can deal with this problem well, but what if we change the iteration to the following:

```
for(int i = 0; i < N; i=i+a[i])
```

Then, the “for” may not work, as the iteration **loop depends on** the values held in the array `a[]`.

Limitation of the For Construct (Continued)

Using the **task** construct to solve it

```
int main(int argc, char* argv[])
{
    int a[N];
    init(a);
    #pragma omp parallel num_threads(2)
    {
        #pragma omp single
        {
            for(int i = 0; i < N; i=i+a[i])
            {
                #pragma omp task
                process(a[i]);
            }
        }
    }
    return 0;
}
```

```
void init(int* a, N)
{
    for(int i=0; i<N; i++)
        a[i] = i;
}

process(a[i]) can be
printf("%d\n", a[i]);
```

Example: Difference between “section” and “task”

```
// sections
#pragma omp sections
{
    #pragma omp section
    foo();
    #pragma omp section
    bar();
}
```

Sections are enclosed within the sections construct and (unless the nowait clause was specified) threads will **not leave** it until all sections have been executed :

```
          [ sections ]
Thread 0: -----< section 1 >----->*-----
Thread 1: -----< section 2          >*-----
Thread 2: ----->*-----
...                                     *
Thread N-1: ----->*-----
```

Example: Difference between “section” and “task”

```
// tasks
#pragma omp single nowait
{
    #pragma omp task
        foo();
    #pragma omp task
        bar();
}
#pragma omp taskwait
//taskwait work like barrier
```

Tasks are queued and executed whenever possible at the so-called task **scheduling points**, which we will talk later.

In fact, If no task scheduling points are present inside the region's code (like above), the OpenMP runtime might start the tasks whenever it deems appropriate. Therefore, **task** is wiser than **sections** while **distributing computing resources** most of the time.

Section v.s. Task (Continued)

Here are possible scenarios of what might happen if there are three threads in last example:

```

      +---+--->[ task queue ]---+
      |   |               |
      |   |               +-----+
      |   |               |
      |   |               |
Thread 0: --< single >-| v |-----
Thread 1: ----->|< foo() >|-----
Thread 2: ----->|< bar() >|-----
    
```

```

      +---+--->[ task queue ]---+
      |   |               |
      |   |               +-----+
      |   |               |
      |   |               |
Thread 0: --< single >-| v |-----
Thread 1: ----->|< foo() >< bar() >|---
Thread 2: ----->|               |---
    
```

```

      +---+--->[ task queue ]---+
      |   |               |
      |   |               +-----+
      |   |               |
      |   |               |
Thread 0: --< single >-| v < bar() >|---
Thread 1: ----->|< foo() >         |---
Thread 2: ----->|               |---
    
```

```

Thread 0: --< single: foo(); bar() >*---
Thread 1: ----->*---
Thread 2: ----->*---
    
```

Data Environment:

- An important consideration for OpenMP programming is the understanding and use of data scoping.
- As OpenMP is based upon the shared memory programming model, **most variables are shared by default.**
- Global variables include:
 - ✓ File scope variables, static
- **But not everything is shared...**
- Private variables include:
 - ✓ Loop index variables
 - ✓ Stack variables in subroutines called from parallel regions

- **OpenMP Task**
 - ✓ What is task
 - ✓ Task generation and execution
 - ✓ **Data environment of tasks**
 - ✓ Task Switching
- **Synchronisation**
 - ✓ Why synchronisation
 - ✓ Synchronisation constructs
 - ✓ Critical and atomic
 - ✓ Flush and lock

- Most variables are shared by default in OpenMP.
- However, the **default for tasks** is usually **firstprivate**, because the task may not be executed until later (and variables may have gone out of scope).
- Variables that are **shared** in all constructs starting from the **innermost** enclosing parallel construct are shared, because the barrier guarantees task completion.

The **task** directive takes the following data attribute clauses that define the data environment of the task:

- **default (shared | none)**
- **private (list)**
- **firstprivate (list)**
- **shared (list)**

shared clause: The **shared** clause declares one or more list items to be shared by tasks generated by a parallel, teams, or task generating construct.

private clause: The **private** clause declares one or more list items to be **private** to a task or to a **SIMD** lane.

firstprivate clause: The **firstprivate** clause declares one or more list items to be **private** to a task, and **initialises** each of them with the value that the corresponding original item has when the construct is encountered.

default clause: Specifies the behavior of **unscoped** variables in a parallel region.

Data Scoping Example

```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;
            // Scope of a:
            // Scope of b:
            // Scope of c:
            // Scope of d:
            // Scope of e:
        }
    }
}
```

Data Scoping Example (Continued)

```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;
            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d: firstprivate
            // Scope of e: private
        }
    }
}
```

The if clause:

- When the if clause argument is **false**, the task is executed **immediately** by the encountering thread.

```
#pragma omp task if(0) //this task is undeferred  
    foo();
```

It's a user directed optimisation

- When the cost of deferring the task is too high, compared to the cost of executing the task code
- To control cache and memory affinity

When/where are tasks complete?

- **At thread barriers, explicit or implicit**
 - ✓ applies to all tasks generated in the current parallel region up to the barrier matches user expectation


- **At task barriers**
i.e. Wait until all tasks defined in the current task have completed.

#pragma omp taskwait

Note: applies only to tasks generated in the current task, not to “descendants” .

The **taskwait** Clause

Example: **taskwait** pseudocode



```
#pragma omp parallel
{
    #pragma omp single
    {
        #pragma omp task A();
        #pragma omp task if (0)
        {
            #pragma omp task B();
            #pragma omp task if (0)
            {
                #pragma omp task C();
                D();
                #pragma omp taskwait
                E();
            }
            #pragma omp taskwait
            F();
        }
        #pragma omp taskwait
        G();
    }
}
```

Task Example: Fibonacci Numbers

Example: Computing Fibonacci Numbers

```
#include <stdio.h>
#include <omp.h>
int fib(int n)
{
    int i, j;
    if (n<2)
        return n;
    else
    {
        #pragma omp task shared(i) firstprivate(n)
        i=fib(n-1);
        #pragma omp task shared(j) firstprivate(n)
        j=fib(n-2);
        #pragma omp taskwait
        return i+j;
    }
}
```

Task scheduling point

```
int main()
{
    int n = 10;
    omp_set_dynamic(0);
    omp_set_num_threads(4);
    #pragma omp parallel shared(n)
    {
        #pragma omp single
        printf ("fib(%d) = %d\n", n, fib(n));
    }
    return 0;
}
```

Task Example: Traverse a Linked List

Example: Parallel iterate through a linked list


```
#pragma omp parallel
{
    #pragma omp for private(p)
    for ( int i =0; i < numlists; i++)
    {
        p = listheads[i] ;
        while (p)
        {
            #pragma omp task firstprivate(p)
            {
                process(p);
            }
            p=next(p);
        }
    }
}
```

Task Example: Tree Traversal

Example: Postorder tree traversal

```
void postorder(node *p)
{
    if (p->left)
#pragma omp task
        postorder(p->left);
    if (p->right)
#pragma omp task
        postorder(p->right);
#pragma omp taskwait // wait for descendants

    process(p->data);
}
```



Task scheduling point

- Parent task suspended until children tasks complete

- **OpenMP Task**
 - ✓ What is task
 - ✓ Task generation and execution
 - ✓ Data environment of tasks
 - ✓ **Task Switching**
- **Synchronisation**
 - ✓ Why synchronisation
 - ✓ Synchronisation constructs
 - ✓ Critical and atomic
 - ✓ Flush and lock

Tied and Untied

- If the code is executed by the same thread from start to finish, the task is **tied**.
- A task is **untied** if the code can be executed by multiple threads, causing different threads to execute different parts of the code.
- **By default, tasks are tied** and can be referred to as untied by using the **untied** clause in conjunction with the task directive.

Task Switching

- Certain constructs have task scheduling points at **defined** locations within them
- When a thread encounters a task scheduling point, it is allowed to suspend the current task and execute another (called **task switching**)
- It can then return to the original task and resume

Where are scheduling points

- The point where task construction is encountered;
- The point where the **taskwait** construct is encountered;
- The point where an implicit or explicit **barrier** is encountered;
- The "mission accomplished" point.

Task Switching Example

- Too many tasks generated in an eye-blink
- Generating task will have to **suspend** for a while

```
#pragma omp single
{
    for (i=0; i<ONEZILLION; i++)
    {
        #pragma omp task
        process(item[i]);
    }
}
```

- With task switching, the executing thread can:
 - ✓ Execute an already generated task (draining the “**task pool**”)
 - ✓ Dive into the encountered task (could be very **cache-friendly**)

Thread Switching

```
#pragma omp single
{
    #pragma omp task untied
    for (i=0; i<ONEZILLION; i++)
    #pragma omp task
        process(item[i]);
}
```

- Eventually, too many tasks are generated
- Generating task is suspended and executing thread switches to a long and boring task
- Other threads get rid of all already generated tasks, and start starving...
- With thread switching, the generating task can be resumed by a different thread, and starvation is over
- The programmer is responsible for synchronisation!


Conclusions on tasks

- Enormous amount of work by many people
- Tightly integrated into 3.0 spec
- Flexible model for irregular parallelism
- Provides balanced solution despite often conflicting goals
- Appears that performance can be reasonable

- **OpenMP Task**
 - ✓ What is task
 - ✓ Task generation and execution
 - ✓ Data environment of tasks
 - ✓ Task Switching
- **Synchronisation**
 - ✓ Why synchronisation
 - ✓ Synchronisation constructs
 - ✓ Critical and atomic
 - ✓ Flush and lock

Race Condition

- A data race occurs when two threads access the same memory without proper **synchronisation**.
- This can cause the program to produce non-deterministic results in parallel mode.



Race Condition!

Race Condition Example

- **Race Condition** is discussed in [Week 3 lecture](#).
- It can be elaborated by following intuitive example

```
THREAD 1:  
update(x)  
{  
    x = x + 1  
}  
  
x = 0  
update(x)  
print(x)
```

```
THREAD 2:  
update(x)  
{  
    x = x + 1  
}  
  
x = 0  
update(x)  
print(x)
```


Race Condition Example (Continued)

- One possible execution sequence:
 - ✓ Thread 1 initialises x to 0 and calls `update(x)`.
 - ✓ Thread 1 adds 1 to x and x now equals 1.
 - ✓ Thread 2 initialises x to 0 and calls `update(x)`. x now equals 0.
 - ✓ Thread 1 prints x, which is equal to 0 instead of 1
 - ✓ Thread 2 adds 1 to x. x now equals 1.
 - ✓ Thread 2 prints x as 1.

```
THREAD 1:  
update(x)  
{  
    x = x + 1  
}  
  
x = 0  
update(x)  
print(x)
```

```
THREAD 2:  
update(x)  
{  
    x = x + 1  
}  
  
x = 0  
update(x)  
print(x)
```

- **OpenMP Task**
 - ✓ What is task
 - ✓ Task generation and execution
 - ✓ Data environment of tasks
 - ✓ Task Switching
- **Synchronisation**
 - ✓ Why synchronisation
 - ✓ **Synchronisation constructs**
 - ✓ Critical and atomic
 - ✓ Flush and lock

Synchronisation Constructs

- To avoid a situation like this, the updating of x must be **synchronised** between the two threads to ensure that the correct result is produced.
- **OpenMP** provides a variety of **Synchronisation Constructs** that control how the execution of each thread proceeds relative to other team threads.
- **OpenMP Synchronisation:**
 - Single/Master ([Week 3 lecture](#))
 - Ordered ([Week 3 lecture](#))
 - Barriers
 - Critical
 - Atomic
 - Flush (memory subsystem synchronisation)
 - Locks

(Lecture 3) Single Worksharing Construct

- The **single** construct denotes a block of code that is executed by only one thread (not necessarily the master thread).
- A barrier is implied at the end of the single block (can remove the barrier with a **nowait** clause).

```
#pragma omp parallel
{
    do_many_things();

#pragma omp single
{
    exchange_boundaries();
}

    do_many_other_things();
}
```

(Lecture 3) Master Construct

- The master construct denotes a structured block that is only executed by the master thread.
- The other threads just skip it (no synchronisation is implied).

```
#pragma omp parallel
{
    do_many_things();
    #pragma omp master
    {
        exchange_boundaries();
    }
    #pragma omp barrier
    do_many_other_things();
}
```

(Lecture 3) Loop Worksharing Constructs:

The **ordered** Clause

The **ordered** region executes in the sequential order.

```
void test(int first, int last)
{
#pragma omp parallel
#pragma omp for schedule(static) ordered
    for (int i = first; i <= last; ++i)
    {
        // Do something here.
        if (i % 2)
        {
            #pragma omp ordered
            printf_s("test() iteration %d\n", i);
        }
    }
}

int main(int argc, char *argv[])
{
    test(1, 8);
}
```

Output:

```
test() iteration 1
test() iteration 3
test() iteration 5
test() iteration 7
```

Exercise:

- Delete “**#pragma omp ordered**”, compile and run the program multiple times.
- Do you see any difference?

(Lecture 3) Loop Worksharing Constructs: The **ordered** Clause (Cont'd)



The **omp ordered** directive must be used as follows:


- It must appear within the extent of an “**omp for**” or “**omp parallel for**” construct containing an ordered clause.
- It applies to the statement block immediately following it. Statements in that block are executed in the same order in which iterations are executed in a sequential loop.
- An iteration of a loop must not execute the same omp ordered directive more than once.
- An iteration of a loop must not execute more than one distinct omp ordered directive.

- **OpenMP Task**
 - ✓ What is task
 - ✓ Task generation and execution
 - ✓ Data environment of tasks
 - ✓ Task Switching
- Synchronisation
 - ✓ Why synchronisation
 - ✓ Synchronisation constructs
 - ✓ **Critical and atomic**
 - ✓ Flush and lock

Synchronisation Constructs

CRITICAL Directive

- **Mutual exclusion**: the **CRITICAL** directive specifies a region of code that must be executed by only one thread at a time.
- If a **thread** is currently executing inside a **CRITICAL** region and another thread reaches that **CRITICAL** region and attempts to execute it, it will block until the first thread exits that **CRITICAL** region.



Usage:

```
#pragma omp critical [ name ] newline
```

- The optional name enables multiple different **CRITICAL** regions to exist:
 - Names act as **global identifiers**. Different **CRITICAL** regions with the same name are treated as the same region.
 - All **CRITICAL** sections which are unnamed, are treated as the same section.

Example:

```
#include <omp.h>
main(int argc, char *argv[]) {
    int x;
    x = 0
    #pragma omp parallel shared(x)
    {
        #pragma omp critical
        x = x + 1;
    }
    /* end of parallel region */
    return 0;
}
```

Notes:

- All threads in the team will attempt to execute in parallel, however, because of the **CRITICAL** construct surrounding the increment of **x**, only one thread will be able to read/increment/write **x** at any time.

CRITICAL Construct/Example

Compare ***a[i]*** and ***max*** again because ***max*** could have been changed by another thread after the comparison outside the critical section.

```
#include <omp.h>
#include <stdio.h>
#include <stdlib.h>
#define SIZE 10
int main() {
    int i, max, a[SIZE];
    for (i = 0; i < SIZE; i++)
    {
        a[i] = rand();
        printf("%d\n", a[i]);
    }
    max = a[0];
    #pragma omp parallel for num_threads(4)
    for (i = 1; i < SIZE; i++)
    {
        if (a[i] > max)
        {
            #pragma omp critical
            if (a[i] > max) {
                max = a[i];
            }
        }
    }
    printf("max = %d\n", max);
    return 0;
}
```

Output:

```
41
18467
6334
26500
19169
15724
11478
29358
26962
24464
max = 29358
```

Notes:

- You might not get the same output because we have used `random` number.

Synchronisation Constructs

ATOMIC Directive

- **ATOMIC** provides **mutual exclusion** but only applies to the load/update of a memory location.
- In essence, this directive provides a **mini-CRITICAL** section.
- The directive applies only to a **single, immediately** following **statement**
- An **ATOMIC** statement must follow a **specific syntax**.
 - **ATOMIC** construct may only be used together with an expression statement with one of operations: **+, *, -, /, &, ^, |, <>**.

ATOMIC Construct/Example

Usage:

```
#pragma omp atomic [ read | write | update | capture ] newline
```

Example:

```
#include <stdio.h>
#include <omp.h>
#define MAX 10
int main() {
    int count = 0;
    #pragma omp parallel num_threads(MAX)
    {
        #pragma omp atomic
        count++;
    }
    printf("Number of threads: %d\n", count);
    return 0;
}
```

Output:

```
Number of threads: 10
```

- **Critical section:**

- Ensures serialisation of blocks of code.
- Can be extended to serialise groups of blocks with proper use of "name" tag.
- Slower!

- **Atomic operation:**

- Is much faster!
- Only ensures the serialisation of a particular operation.

- **OpenMP Task**
 - ✓ What is task
 - ✓ Task generation and execution
 - ✓ Data environment of tasks
 - ✓ Task Switching
- **Synchronisation**
 - ✓ Why synchronisation
 - ✓ Synchronisation constructs
 - ✓ Critical and atomic
 - ✓ **Flush and lock**

Synchronisation Constructs

FLUSH Directive

- The **FLUSH** directive identifies a **synchronisation** point at which the implementation must provide a consistent view of memory. **Thread-visible** variables are written back to memory at this point.
- Usage: `#pragma omp flush (list) newline`
- The optional **list** contains a list of named variables that will be flushed in order to avoid flushing all variables.
- For **pointers** in the list, note that the pointer itself is flushed, not the object it points to.

Synchronisation Constructs

FLUSH Directive

- Implementations must ensure any prior modifications to **thread-visible** variables are visible to all threads after this point;
 - ie. compilers must restore values from registers to memory, hardware might need to flush write buffers, etc.
- The **FLUSH** directive is **implied** for the directives shown in the table below. The directive is not implied if a **NOWAIT** clause is present.

C / C++

barrier
parallel - upon entry and exit
critical - upon entry and exit
ordered - upon entry and exit
for - upon exit
sections - upon exit
single - upon exit

Synchronisation Constructs

FLUSH Directive

```
#include <stdio.h>
#include <omp.h>
void read(int *data) {
    printf("read data\n");
    *data = 1;
}
void process(int *data) {
    printf("process data\n");
    (*data)++;
}
int main()
{
    int data;
    int flag;
    flag = 0;
    #pragma omp parallel sections num_threads(2)
    {
        #pragma omp section
        {
            printf("Thread %d: ", omp_get_thread_num( ));
            read(&data);
            #pragma omp flush(data)
            flag = 1;
            #pragma omp flush(flag)
            // Do more work.
        }
    }
}
```

```
#pragma omp section
{
    while (!flag) {
        #pragma omp flush(flag)
    }
    #pragma omp flush(data)

    printf("Thread %d: ", omp_get_thread_num());
    process(&data);
    printf("data = %d\n", data);
}
}
return 0;
}
```

Output:

```
Thread 0: read data
Thread 1: process data
data = 2
```

Synchronisation Constructs

Lock Directive

- A **lock** implies a memory fence of all thread visible variables.
- These routines are used to guarantee that only one thread accesses a variable at a time to avoid **race conditions**.
- C/C++ lock variables must have type “**omp_lock_t**” or “**omp_nest_lock_t**”.
- All lock functions require an argument that has a pointer to **omp_lock_t** or **omp_nest_lock_t**.
- Simple Lock routines:
 - `omp_init_lock(omp_lock_t*);`
 - `omp_set_lock(omp_lock_t*);`
 - `omp_unset_lock(omp_lock_t*);`
 - `omp_test_lock(omp_lock_t*);`
 - `omp_destroy_lock(omp_lock_t*);`

General Procedure to Use Locks

- Define the **lock** variables.
- Initialise the **lock** via a call to **omp_init_lock**.
- Set the **lock** using **omp_set_lock** or **omp_test_lock**. The latter checks whether the **lock** is actually available before attempting to set it. It is useful to achieve asynchronous thread execution
- Unset a **lock** after the work is done via a call to **omp_unset_lock**.
- Remove the **lock** association via a call to **omp_destroy_lock**.

Synchronisation Constructs

Lock Directive

Example:

```
#include <stdlib.h>
#include <stdio.h>
#include <omp.h>

int main()
{
    int x;
    omp_lock_t lck;

    omp_init_lock (&lck);
    omp_set_lock (&lck);
    x = 0;

    #pragma omp parallel shared (x)
    {
        #pragma omp master
        {
            x = x + 1;
            omp_unset_lock (&lck);
        }
    }
    omp_destroy_lock (&lck);
    return 0;
}
```

- Readings
 - [Section vs Task](#)
 - [OpenMP Tasking Tutorial](#)
 - [Lock vs nested lock](#)



Copyright Notice

Material used in this recording may have been reproduced and communicated to you by or on behalf of **The University of Western Australia** in accordance with section 113P of the *Copyright Act 1968*.

Unless stated otherwise, all teaching and learning materials provided to you by the University are protected under the Copyright Act and is for your personal use only. This material must not be shared or distributed without the permission of the University and the copyright owner/s.